

Package ‘spdm’

February 3, 2019

Type Package

Title Functions for working with symmetric positive-definite matrices

Version 1.0.0

Author Corson N. Areshenkoff

Maintainer Corson N. Areshenkoff <areshenk@protonmail.com>

Description Means, estimation, and operations for symmetric positive-definite matrices.

Depends R (>= 3.0.2),

Imports expm, huge, nlshrink, kernlab, geigen

License GPL-3

Encoding UTF-8

LazyData true

RoxygenNote 6.1.0

Suggests knitr,
rmarkdown

VignetteBuilder knitr

R topics documented:

is.spd	2
spd.correlation	2
spd.dist	3
spd.estimate	3
spd.expmap	4
spd.interpolate	5
spd.logmap	5
spd.mean	6
spd.pca	7
spd.pca.predict	7
spd.transport	8
spd.vectorize	8
spd.whiten	9

Index	10
--------------	-----------

is.spd

Check if a matrix is symmetric and positive definite

Description

Check if a matrix is symmetric and positive definite

Usage

```
is.spd(..., tol = 1e-08)
```

Arguments

...	One or more numeric matrices
tol	Numerical tolerance for checking symmetry and positive definiteness

Details

Function checks whether the matrix *x* is symmetric and positive definite. Symmetry is evaluated up to an entrywise tolerance of *tol*, so that differences smaller than *tol* are ignored. Positive-definiteness is checked by computing the eigenvalues of *x* using `eigen`, setting eigenvalues smaller than *tol* in absolute value to zero, and then checking whether any are less than or equal to zero.

Value

A Boolean value indicating whether the matrix is symmetric and positive definite

spd.correlation

Compute (partial) correlations

Description

Transforms an SPD matrix into a matrix of partial correlations

Usage

```
spd.correlation(x, method = "correlation")
```

Arguments

<i>x</i>	An SPD matrix to be whitened
<i>method</i>	A string specifying either "correlation" (default) or "partial" correlation.

Value

A symmetric, positive-definite matrix.

spd.dist	<i>Distance between two symmetric, positive-definite matrices</i>
----------	---

Description

Function implements several distance measures between symmetric, positive-definite matrices.

Usage

```
spd.dist(x, y, method = "euclidean", ...)
```

Arguments

x, y	Symmetric, positive-definite matrices
method	The distance measure. See details.

Details

Allowable distance measures are

- "euclidean": The Frobenius norm of the difference $x-y$.
- "logeuclidean": The Frobenius norm of the difference $\log(x)-\log(y)$ in the tangent space..
- "cholesky": The Frobenius norm of the difference between the cholesky factors of x and y . Not affinely invariant.
- "riemannian": The Riemmanian distance proposed by Barachant, et al. (2013)
- "stein": The square root of Jensen;Bregman Log determinant divergence

spd.estimate	<i>Covariance estimation</i>
--------------	------------------------------

Description

Function implements several forms of covariance estimation.

Usage

```
spd.estimate(x, method = "linshrink", ...)
```

Arguments

x	A data matrix, where rows are observations and columns are variables
method	Method of covariance estimation. See details
...	Additional arguments passed to estimation functions. See details.

Details

Allowable estimation methods are:

- "sample": The ordinary sample covariance. Generally a poor choice in anything but very low dimensional settings, and is not guaranteed to be positive-definite.
- "linshrink": Linear shrinkage estimator proposed by Ledoit and Wolf (2004)
- "nlshrink": Non-linear shrinkage estimator proposed by Ledoit and Wolf (2012)
- "glasso": Graphical lasso (glasso) estimation using the huge package. Typically generates sparse estimates.

Additional arguments may be passed to the functions which perform estimation. Specifically:

- "sample": Uses `cov(x, ...)`
- "linshrink": Uses `nlshrink::linshrink_cov(x, ...)`
- "nlshrink": Uses `nlshrink::nlshrink_cov(x, ...)`
- "glasso": Uses `huge(x, method = 'glasso', cov.output = T, ...)` followed by `huge.select(x, ...)`. Note that method cannot be overridden, as other estimation methods do not return covariance estimates. Additional arguments to `huge()` or `huge.select()` should be prepended with `huge.` or `select.`, respectively.

In all cases, function will generate a warning if the estimated matrix is not positive definite.

Value

A covariance matrix

spd.expmap

Projection from the tangent space

Description

Function projects a tangent vector x (a symmetric matrix) at a point p onto the space of SPD matrices.

Usage

```
spd.expmap(x, p = NULL)
```

Arguments

x	A symmetric matrix
p	The point to whose tangent space x belongs

Value

A symmetric, positive-definite matrix.

spd.interpolate	<i>Interpolation between two symmetric, positive-definite matrices</i>
-----------------	--

Description

Function smoothly interpolates between two SPD matrices x and y . The interpolation is parametrized by t , where $t=0$ returns x , $t=1$ returns y , and $t=.5$ returns the mean of x and y

Usage

```
spd.interpolate(x, y, t, method = "euclidean", ...)
```

Arguments

x , y	Symmetric, positive-definite matrices
t	Interpolation parameter in $[0, \text{infinity})$
method	Type of interpolation (see details)

Details

Allowable distance measures are

- "euclidean": Euclidean interpolation $(1-t)x + ty$
- "logeuclidean": Euclidean interpolation in the tangent space.
- "riemannian": Interpolation along the geodesic path from x to y

spd.logmap	<i>Projection onto the tangent space</i>
------------	--

Description

Function projects an SPD matrix x onto the tangent space at a point p

Usage

```
spd.logmap(x, p = NULL)
```

Arguments

x	A symmetric positive definite matrix
p	The point on whose tangent space to project x

Value

A symmetric matrix.

spd.mean

Compute the mean of a set of spd matrices

Description

Function computes the mean of a set of symmetric positive-definite matrices. Several methods are implemented, as described in Details.

Usage

```
spd.mean(x, method = "euclidean", ...)
```

Arguments

x	A list of symmetric, positive-definite matrices
method	The type of mean to compute. See details
...	Further arguments. See details

Details

Function computes the mean of a set of symmetric, positive-definite matrices. Several methods are implemented:

- "euclidean": The ordinary arithmetic mean – the sum of the matrices in x, divided by the number of matrices. This is guaranteed to be an spd matrix, but does not necessarily preserve the spectral characteristics of the individual matrices.
- "logeuclidean": Computed by taking the arithmetic mean of the logarithms of the matrices in x, and then projecting back onto the space of spd matrices. In general, better behaved than the arithmetic mean.
- "riemannian": The Riemmanian p-mean. Smoothly interpolates between the harmonic mean at $p = -1$ to the geometric mean at $p = 0$ to the arithmetic mean at $p = 1$. Is approximated iteratively using the fixed point algorithm described by Congedo, Barachant and Koopaei (2017). Requires a parameter p in the interval $[-1, 1]$ specifying the type of mean, a maximum error tolerance tol (default .01), and a maximum number of iterations (default 50). The case $p = 0$ is approximated by computing the p-means at -1 and 1 , and then returning the midpoint between them using `spd.interpolate(..., method = 'riemannian')`

Value

The mean of the matrices in x.

spd.pca	<i>Kernel pca for SPD matrices</i>
---------	------------------------------------

Description

Function performs kernel principal component analysis on a set of symmetric, positive-definite matrices using an rbf kernel: $\exp(-\text{sigma} * d(i,j)^2)$, where $d(i,j)$ is a distance function implemented by `spd.dist`. This function is more or less a wrapper around the kernlab function `kpca`.

Usage

```
spd.pca(x, method = "euclidean", sigma = 1, ...)
```

Arguments

<code>x</code>	A list of symmetric, positive-definite matrices
<code>method</code>	The type of distance. See <code>spd.dist</code> , but also see details. Defaults to "logeuclidean".
<code>sigma</code>	scale parameter for rbf kernel
<code>...</code>	Further arguments for <code>kernlab::kpca</code> .

Details

Function performs `kpca` using a rbf kernel, where the distance between two inputs is given by `method`. Note that only "euclidean" and "logeuclidean" have been proven to give rise to positive-definite kernels for all values of `sigma`, although any distance implemented in `spd.dist` may be used. Anecdotally, `method = "riemannian"` often achieves superior performance.

Value

An S4 object of class `kpca`.

spd.pca.predict	<i>Prediction for SPD kernel PCA</i>
-----------------	--------------------------------------

Description

Function accepts a fitted `kpca` object returned by `spd.pca` and a list of SPD matrices and returns a matrix of estimated principal component scores.

Usage

```
spd.pca.predict(fit, x)
```

Arguments

<code>fit</code>	An object of class <code>kpca</code> return by <code>spd.pca</code>
<code>x</code>	A list of SPD matrices for which to derive component scores.

Value

A matrix of principal component scores

<code>spd.transport</code>	<i>Parallel transport of a tangent vector</i>
----------------------------	---

Description

Parallel transports a vector `x` in the tangent space at "`from`" to the tangent space at "`to`" using the Schild's ladder algorithm.

Usage

```
spd.transport(x, to, from, nsteps = 10)
```

Arguments

<code>x</code>	A tangent vector (symmetric matrix)
<code>to</code>	The SPD matrix to whose tangent space to move <code>x</code>
<code>from</code>	The SPD matrix to whose tangent space <code>x</code> belongs
<code>nsteps</code>	The number of steps in the geodesic connecting <code>to</code> and <code>from</code>

Value

A symmetric matrix – a tangent vector at `to`.

<code>spd.vectorize</code>	<i>Vectorize and unvectorize a matrix</i>
----------------------------	---

Description

Function converts between a square symmetric matrix and a vector of the lower triangular elements + diagonal. Allows optional scaling of the off-diagonal entries in order to preserve the norm.

Usage

```
spd.vectorize(x, scaling = F)
```


Arguments

x	Either a square matrix, or a vector contains the lower triangular elements of x (including the diagonal elements)
scaling	If TRUE, scales the off diagonal elements by sqrt(2) to preserve the norm of the resulting vector

Details

If input is a square matrix, converts the lower triangular elements (including the diagonal) into a numeric vector. If input is a numeric vector, converts the input to a symmetric matrix. Note that, if the input is a vector, its length must be a triangular number.

spd.whiten	<i>Covariance whitening transform</i>
------------	---------------------------------------

Description

Whitens an SPD matrix using the procedure advocated by Ng, et al. (2014)

Usage

```
spd.whiten(x, p, unwhiten = F)
```

Arguments

x	An SPD matrix to be whitened
p	The SPD baseline whitening matrix
unwhiten	Logical. If TRUE, reverses a previously applied whitening transform.

Value

A symmetric, positive-definite matrix.

Index

`is.spd`, [2](#)

`spd.correlation`, [2](#)

`spd.dist`, [3](#)

`spd.estimate`, [3](#)

`spd.expmap`, [4](#)

`spd.interpolate`, [5](#)

`spd.logmap`, [5](#)

`spd.mean`, [6](#)

`spd.pca`, [7](#)

`spd.pca.predict`, [7](#)

`spd.transport`, [8](#)

`spd.vectorize`, [8](#)

`spd.whiten`, [9](#)